

Hands On Lab: MySQL Overview & Architecture

LAB 1.1: Installing MySQL on a Linux VM

Objective: Install MySQL Server on a Ubuntu 20.04 LTS VM and verify the installation.

Prerequisites:

- Ubuntu 20.04 LTS VM with sudo access
- At least 2GB RAM, 20GB disk space
- Network connectivity

Step-by-Step Instructions:

1. Update System Packages

```
sudo apt update
```

```
sudo apt upgrade -y
```

2. Install MySQL Server

```
sudo apt install -y mysql-server mysql-client
```

3. Verify Installation

```
mysql --version
```

```
mysqld --version
```

4. Start MySQL Service

```
sudo systemctl start mysql
```

```
sudo systemctl enable mysql
```

```
sudo systemctl status mysql
```

5. Connect to MySQL

```
sudo mysql -u root
```

6. Run Initial Setup Commands

```
SELECT VERSION();
```

```
SHOW VARIABLES LIKE 'datadir';
```

```
SHOW VARIABLES LIKE 'basedir';
```

```
SELECT USER();
```

7. Secure MySQL Installation

```
sudo mysql_secure_installation
```

- Set root password (if you get error in this step, you can see the resolution in Appendix A)
- Remove anonymous users
- Disable remote root login
- Remove test database
- Reload privilege tables

8. Verify Security Settings

```
SELECT host, user FROM mysql.user;
```

Hands-on Tasks:

- Document the MySQL version and installation path
- Take screenshots of key directories: /var/lib/mysql, /etc/mysql
- List all MySQL processes running: ps aux | grep mysql
- Check MySQL port: sudo netstat -tuln | grep 3306

Troubleshooting:

- If service fails to start: Check /var/log/mysql/error.log
- Permission issues: Verify /var/lib/mysql ownership: ls -la /var/lib/mysql

LAB 1.2: Exploring Storage Engines

Objective: Understand different MySQL storage engines and their characteristics.

Step-by-Step Instructions:

1. Connect to MySQL

```
mysql -u root -p
```

2. List Available Storage Engines

```
SHOW ENGINES;
```

```
SHOW ENGINES\G
```

3. Create Sample Databases

```
CREATE DATABASE storage_engines_demo;
```

```
USE storage_engines_demo;
```

4. Create Tables with Different Storage Engines

a) InnoDB Table (Default)

```
CREATE TABLE innodb_table (
    id INT PRIMARY KEY AUTO_INCREMENT,
    name VARCHAR(100),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
) ENGINE=InnoDB;
```

b) MyISAM Table

```
CREATE TABLE myisam_table (
    id INT PRIMARY KEY AUTO_INCREMENT,
    description TEXT
) ENGINE=MyISAM;
```

c) Memory Table

```
CREATE TABLE memory_table (
    id INT PRIMARY KEY AUTO_INCREMENT,
    session_data VARCHAR(255)
) ENGINE=MEMORY;
```

d) Archive Table

```
CREATE TABLE archive_table (
    id INT PRIMARY KEY AUTO_INCREMENT,
    log_data TEXT
) ENGINE=Archive;
```

5. Insert Test Data

```
INSERT INTO innodb_table (name) VALUES
('John'), ('Alice'), ('Bob');
```

```
INSERT INTO myisam_table (description) VALUES
('Test data 1'), ('Test data 2');
```

```
INSERT INTO memory_table (session_data) VALUES
('Session 1'), ('Session 2');
```

```
INSERT INTO archive_table (log_data) VALUES
('Log entry 1'), ('Log entry 2');
```

6. Query Information Schema to View Engine Details

```
SELECT TABLE_SCHEMA, TABLE_NAME, ENGINE, TABLE_TYPE
FROM INFORMATION_SCHEMA.TABLES
WHERE TABLE_SCHEMA = 'storage_engines_demo';
```

7. Compare Engine Characteristics

```
SELECT
    t.TABLE_NAME,
    t.ENGINE,
    t.ROW_FORMAT,
    t.TABLE_ROWS,
    ROUND(((t.DATA_LENGTH + t.INDEX_LENGTH) / 1024 / 1024), 2) AS size_mb
```

```
FROM INFORMATION_SCHEMA.TABLES t  
WHERE t.TABLE_SCHEMA = 'storage_engines_demo';
```

Hands-on Tasks:

- Create tables with each engine and observe storage characteristics
- Insert sample data (100+ rows) into each table
- Compare table sizes using INFORMATION_SCHEMA
- Document the SHOW CREATE TABLE output for each engine
- Identify which engines support transactions

Solution to Hands-on Tasks:

Create one table per engine, insert sample data, then compare their on-disk size and behavior using SHOW TABLE STATUS or INFORMATION_SCHEMA.TABLES.

Step 1: Create demo tables

Use a test database so the objects are isolated.

```
CREATE DATABASE IF NOT EXISTS engine_demo;  
USE engine_demo;
```

Create the same logical table with different engines.

```
CREATE TABLE customer_innodb (  
    id INT PRIMARY KEY AUTO_INCREMENT,  
    name VARCHAR(100),  
    email VARCHAR(200),  
    created_at DATETIME  
) ENGINE = InnoDB;
```

```
CREATE TABLE customer_myisam (  
    id INT PRIMARY KEY AUTO_INCREMENT,  
    name VARCHAR(100),  
    email VARCHAR(200),
```

```
created_at DATETIME  
 ) ENGINE = MyISAM;  
  
CREATE TABLE customer_memory (  
    id INT PRIMARY KEY AUTO_INCREMENT,  
    name VARCHAR(100),  
    email VARCHAR(200),  
    created_at DATETIME  
 ) ENGINE = MEMORY;
```

```
CREATE TABLE customer_archive (  
    id INT PRIMARY KEY AUTO_INCREMENT,  
    name VARCHAR(100),  
    email VARCHAR(200),  
    created_at DATETIME  
 ) ENGINE = ARCHIVE;
```

Step 2: Insert identical data

Populate each table with the same rows to make comparisons meaningful.

```
INSERT INTO customer_innodb (name, email, created_at)  
SELECT CONCAT('User', n),  
       CONCAT('user', n, '@example.com'),  
       NOW()  
FROM (  
    SELECT 1 AS n UNION ALL SELECT 2 UNION ALL SELECT 3 UNION ALL SELECT 4 UNION ALL  
    SELECT 5  
    UNION ALL SELECT 6 UNION ALL SELECT 7 UNION ALL SELECT 8 UNION ALL SELECT 9 UNION  
    ALL SELECT 10  
) AS t;
```

```
INSERT INTO customer_myisam (name, email, created_at)  
SELECT name, email, created_at FROM customer_innodb;
```

```
INSERT INTO customer_memory (name, email, created_at)  
SELECT name, email, created_at FROM customer_innodb;
```

```
INSERT INTO customer_archive (name, email, created_at)  
SELECT name, email, created_at FROM customer_innodb;
```

Use one base INSERT and then copy that data into the other engine tables so all four have identical rows.

1. Insert into the InnoDB table

```
USE engine_demo;
```

```
INSERT INTO customer_innodb (name, email, created_at)
```

```
VALUES
```

```
('Alice', 'alice@example.com', NOW()),  
('Bob', 'bob@example.com', NOW()),  
('Charlie','charlie@example.com',NOW()),  
('David', 'david@example.com', NOW()),  
('Emma', 'emma@example.com', NOW()),  
('Frank', 'frank@example.com', NOW()),  
('Grace', 'grace@example.com', NOW()),  
('Helen', 'helen@example.com', NOW()),  
('Ivan', 'ivan@example.com', NOW()),  
('Julia', 'julia@example.com', NOW());
```

This uses a single INSERT with multiple value lists, which is the recommended pattern for inserting many rows efficiently.

2. Copy the same data to other engines

Now clone rows from customer_innodb into the MyISAM, MEMORY, and ARCHIVE tables with INSERT ... SELECT.

```
INSERT INTO customer_myisam (name, email, created_at)
```

```
SELECT name, email, created_at
```

```
FROM customer_innodb;
```

```
INSERT INTO customer_memory (name, email, created_at)
```

```
SELECT name, email, created_at
```

```
FROM customer_innodb;
```

```
INSERT INTO customer_archive (name, email, created_at)
```

```
SELECT name, email, created_at
```

```
FROM customer_innodb;
```

INSERT INTO ... SELECT copies column values directly from the source table and is a standard way to bulk-populate another table with the same structure.

3. Quick verification

```
SELECT 'innodb' AS engine_name, COUNT(*) AS rows FROM customer_innodb
```

```
UNION ALL
```

```
SELECT 'myisam', COUNT(*) FROM customer_myisam
```

```
UNION ALL
```

```
SELECT 'memory', COUNT(*) FROM customer_memory
```

```
UNION ALL
```

```
SELECT 'archive', COUNT(*) FROM customer_archive;
```

All four tables should show the same row count (10 rows in this example), ensuring a fair comparison of storage characteristics.

For a stronger effect, scale this up to a few thousand rows using a numbers helper or stored procedure.

Step 3: Observe storage with SHOW TABLE STATUS

Check per-table size and characteristics.

```
SHOW TABLE STATUS FROM engine_demo LIKE 'customer_%'\G
```

Key columns to discuss with participants:

- Engine: Confirms storage engine used.
- Rows: Estimated number of rows (MyISAM is often exact; InnoDB is approximate).
- Data_length:
 - MyISAM: size of the data file in bytes.
 - InnoDB: approximate space allocated for the clustered index (data) in bytes.
- Index_length:
 - MyISAM: size of the index file.
 - InnoDB: approximate space for secondary indexes.
- Row_format, Data_free: Observe fragmentation and free space especially for InnoDB/MyISAM.

Prompt questions for learners:

- Compare Data_length and Index_length between InnoDB and MyISAM for the same row count.
- Check how MEMORY table reports size (entirely in RAM, still exposed via metadata).
- Note ARCHIVE's compact storage and append-only nature (no index_length by default).

Step 4: Observe via INFORMATION_SCHEMA.TABLES

Use a single query to get sizes in MB for each engine.

```
SELECT
```

```
    table_name,
```

```
    engine,
```

```

table_rows,
ROUND(data_length / 1024 / 1024, 2) AS data_mb,
ROUND(index_length / 1024 / 1024, 2) AS index_mb,
ROUND((data_length + index_length) / 1024 / 1024, 2) AS total_mb

FROM information_schema.tables

WHERE table_schema = 'engine_demo'

AND table_name LIKE 'customer_%'

ORDER BY total_mb DESC;

```

Discuss with the group:

- Which engine uses more total space for the same data set and why (e.g., InnoDB overhead vs MyISAM)?
- MEMORY table will show size but is volatile and lost on restart.
- ARCHIVE table typically has small index_length but may have more compressed data_length for large, append-only workloads.

Step 5: Optional extensions

To deepen the lab:

- Add indexes only on some engines and re-run the size query to show impact of INDEX_LENGTH.
- Delete a large percentage of rows from the InnoDB/MyISAM tables, then re-check Data_length and Data_free to show fragmentation and need for OPTIMIZE TABLE or ALTER TABLE for space reclamation.
- Restart MySQL and show that MEMORY table data is gone while others persist, tying storage characteristics to durability guarantees.

Assessment Questions:

1. Which engine is the default in MySQL 8.0?
2. Which engine stores data in memory?
3. Which engine is read-only but compresses well?
4. Which engine supports foreign keys?

APPENDIX - A

1. ... Failed! Error: SET PASSWORD has no significance for user 'root'@'localhost' as the authentication method used doesn't store authentication data in the MySQL server. Please consider using ALTER USER instead if you want to change authentication parameters.

This error is **expected behavior in newer MySQL versions** and is not a failure in the usual sense 

Let me explain **why it happens** and **how to fix it correctly**.

Why this error occurs

Your MySQL **root user** is authenticated using an **external authentication plugin**, most commonly:

- auth_socket (Linux / Ubuntu)
- caching_sha2_password managed externally

When this is the case:

- SET PASSWORD **does nothing**
- MySQL does **not store the password internally**
- Hence the warning/error:

SET PASSWORD has no significance... authentication method doesn't store authentication data

This is very common in:

- MySQL 8.x
 - Ubuntu/Debian installations
 - Docker-based MySQL setups
-

Correct ways to resolve it

Option 1: Change password using ALTER USER (Recommended)

If you want MySQL to manage the password internally:

```
ALTER USER 'root'@'localhost'  
IDENTIFIED WITH mysql_native_password  
BY 'NewStrongPassword!';
```

Then refresh privileges:

```
FLUSH PRIVILEGES;
```

- ✓ This **switches authentication** to a password-based mechanism
 - ✓ Fixes the error permanently
-

Option 2: Keep auth_socket (Passwordless root login)

If you want root to log in **only via OS user** (common in servers):

```
sudo mysql
```

Check authentication method:

```
SELECT user, host, plugin FROM mysql.user WHERE user='root';
```

If plugin = auth_socket, then:

- Password changes are irrelevant
 - This is **secure and valid**
 - No action needed
-

Option 3: Create a separate admin user (Best practice)

Recommended for production systems

```
CREATE USER 'dbadmin'@'localhost'  
IDENTIFIED BY 'StrongPassword!';
```

```
GRANT ALL PRIVILEGES ON *.* TO 'dbadmin'@'localhost'  
WITH GRANT OPTION;
```

```
FLUSH PRIVILEGES;
```

Use dbadmin instead of root.

 **Which option should you choose?**

Scenario	Best Option
Local development	Option 1
Production server	Option 3
Linux-only admin access	Option 2
Docker / CI pipelines	Option 1 or 3

 **Verify the fix**

`mysql -u root -p`

or

`SHOW GRANTS FOR 'root'@'localhost';`

 **Important Note**

 **Avoid using root with passwords in production**

Create dedicated users with least privileges instead.