

InnoDB Engine Deep Dive

3.1 ACID Properties

What are ACID Properties?

ACID is a set of database properties that ensure reliable transactions. Only InnoDB (and similar ACID-compliant engines) guarantees these properties.

Atomicity (A)

Definition: Transaction either completes fully or not at all

Guarantee: All statements in a transaction execute together

BEGIN;

UPDATE accounts SET balance = balance - 100 WHERE id = 1;

UPDATE accounts SET balance = balance + 100 WHERE id = 2;

COMMIT; -- Both updates succeed or both fail (no partial update)

Example of Atomicity Protection:

-- Without atomicity (problematic)

-- If server crashes after first UPDATE, data is inconsistent

-- With ACID (InnoDB)

-- Either both succeed or both fail

BEGIN;

INSERT INTO orders (customer_id, total) VALUES (1, 500);

INSERT INTO order_items (order_id, product_id) VALUES (LAST_INSERT_ID(), 42);

COMMIT; -- If ROLLBACK occurs, both INSERTs are undone

Consistency (C)

Definition: Database moves from one valid state to another valid state

Constraints Maintained:

- Primary key uniqueness

- Foreign key relationships
- Column data types and constraints

-- *Violation attempt*

INSERT INTO users (email) VALUES (NULL); -- *If NOT NULL constraint exists*

-- *Foreign key violation*

INSERT INTO orders (customer_id) VALUES (999); -- *If customer 999 doesn't exist*

-- *Error: Cannot add or update a child row*

-- *Atomicity ensures consistency:*

-- *If transaction violates constraint, entire transaction rolls back*

Isolation (I)

Definition: Concurrent transactions don't interfere with each other

Isolation Levels (increasing isolation):

1. READ UNCOMMITTED

SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;

-- *Dirty reads possible: can see uncommitted data*

-- *Risk: Other transaction may rollback*

2. READ COMMITTED (Most common in production)

SET TRANSACTION ISOLATION LEVEL READ COMMITTED;

-- *Only committed data visible*

-- *Non-repeatable reads possible: data may change within transaction*

3. REPEATABLE READ (MySQL default)

SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;

-- *Data read is guaranteed not to change during transaction*

-- *Phantom reads possible: new rows may appear matching WHERE clause*

4. SERIALIZABLE

SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;

-- *Highest isolation, slowest performance*

-- *Transactions execute as if sequential*

Isolation Levels Comparison:

Isolation Level	Dirty Read	Non-repeatable Read	Phantom Read
READ UNCOMMITTED	Yes	Yes	Yes
READ COMMITTED	No	Yes	Yes
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No

Durability (D)

Definition: Committed data persists despite failures

Guarantees:

- Data written to persistent storage
- Survives power failures, crashes
- Recovered on restart

-- *Once COMMIT succeeds, data is durable*

BEGIN;

INSERT INTO users (name) VALUES ('John');

COMMIT; -- *Data is now permanent on disk*

-- *Even if server crashes immediately after COMMIT,*

-- *data survives and is recovered on restart*

Durability Mechanisms:

- Write-Ahead Logging (WAL): Log written before data
- Redo log: Replays committed transactions on recovery
- InnoDB recovery process: Automatic recovery on startup

3.2 InnoDB Transactions

Transaction Basics

Starting a Transaction:

BEGIN; -- *Start transaction*

START TRANSACTION; -- *Alternative*

-- *Multiple statements within transaction*

SELECT * FROM accounts WHERE id = 1 FOR UPDATE;

INSERT INTO log (action) VALUES ('Transfer started');

UPDATE accounts SET balance = balance - 100 WHERE id = 1;

UPDATE accounts SET balance = balance + 100 WHERE id = 2;

INSERT INTO transactions (from_id, to_id, amount)

VALUES (1, 2, 100);

COMMIT; -- *Save all changes*

COMMIT vs ROLLBACK

COMMIT: Saves all changes permanently

BEGIN;

UPDATE inventory SET quantity = quantity - 1 WHERE product_id = 42;

SELECT quantity FROM inventory WHERE product_id = 42; -- *Shows new value*

COMMIT; -- *Changes permanently saved*

ROLLBACK: Discards all changes

BEGIN;

UPDATE inventory SET quantity = quantity - 1 WHERE product_id = 42;

ROLLBACK; -- *Changes discarded, reverts to state before BEGIN*

SELECT quantity FROM inventory WHERE product_id = 42; -- *Shows original value*

Savepoints

Savepoints allow partial rollback within a transaction:

```
BEGIN;
```

```
UPDATE accounts SET balance = balance - 100 WHERE id = 1;
```

```
SAVEPOINT sp1; -- Create savepoint
```

```
UPDATE accounts SET balance = balance + 100 WHERE id = 2;
```

-- Error occurs!

```
ROLLBACK TO sp1; -- Rollback only second update
```

-- First update still in transaction

```
UPDATE accounts SET balance = balance + 100 WHERE id = 3;
```

```
COMMIT; -- Saves first and third updates
```

Autocommit Mode

By default, MySQL runs in autocommit mode:

-- Check autocommit status

```
SHOW VARIABLES LIKE 'autocommit'; -- Default: ON
```

-- Disable autocommit for explicit transactions

```
SET autocommit = 0;
```

-- Now each statement must be explicitly committed

-- Re-enable

```
SET autocommit = 1;
```

-- Or use explicit BEGIN to disable autocommit for one transaction

BEGIN;

-- Even with autocommit ON, transaction requires explicit COMMIT

COMMIT;

3.3 Locking and Concurrency

Row-Level Locking

InnoDB uses row-level locks for fine-grained concurrency control:

-- Shared lock (read lock)

```
SELECT * FROM accounts WHERE id = 1 LOCK IN SHARE MODE;
```

-- Exclusive lock (write lock) - used in UPDATE/DELETE

```
SELECT * FROM accounts WHERE id = 1 FOR UPDATE;
```

-- These locks prevent other transactions from modifying/reading the row

Lock Conflicts Example

Transaction 1:

```
BEGIN;
```

```
SELECT * FROM accounts WHERE id = 1 FOR UPDATE; -- Acquires exclusive lock
```

-- Lock held until COMMIT

Transaction 2 (in parallel):

```
BEGIN;
```

```
UPDATE accounts SET balance = 500 WHERE id = 1; -- Waits for lock
```

-- Blocked until Transaction 1 releases lock

Lock Monitoring

-- View current locks

```
SHOW OPEN TABLES WHERE In_use > 0;
```

-- View lock waits (MySQL 8.0+)

```
SELECT * FROM INFORMATION_SCHEMA.INNODB_LOCKS;
```

```
SELECT * FROM INFORMATION_SCHEMA.INNODB_LOCK_WAITS;
```

```
-- Kill blocking transaction
```

```
SHOW PROCESSLIST;
```

```
KILL QUERY process_id;
```

```
KILL CONNECTION process_id;
```

Gap Locks and Next-Key Locks

Gap Lock: Prevents insertion of new rows in a range

-- Range query locks gap between rows

```
SELECT * FROM inventory
```

```
WHERE quantity < 10
```

```
FOR UPDATE;
```

-- Gap locks prevent new rows with quantity < 10 from being inserted

Next-Key Lock: Combination of record and gap lock

3.4 InnoDB Tablespaces

What is a Tablespace?

Tablespace is the storage unit in InnoDB that contains table data, indexes, and metadata.

Tablespace Types

1. System Tablespace (mysql)

- Stores data dictionary
- Undo space
- Change buffer
- Configuration: innodb_data_file_path

2. File-Per-Table Tablespaces

[mysqld]

innodb_file_per_table = ON *# Default in MySQL 5.7+*

Each table gets its own .ibd file:

```
ls -lh /var/lib/mysql/database_name/  
# table1.ibd table2.ibd table3.ibd
```

3. General Tablespaces

-- Create general tablespace

```
CREATE TABLESPACE ts1 ADD DATAFILE '/var/lib/mysql/ts1.ibd';
```

-- Create table in specific tablespace

```
CREATE TABLE inventory (  
    id INT PRIMARY KEY,  
    product_name VARCHAR(255)  
) TABLESPACE ts1;
```

4. Undo Tablespaces

- Stores undo logs for transaction rollback

- Separate from data tablespaces
- Configuration: innodb_undo_tablespaces

Creating and Managing Tablespaces

-- *List tablespaces*

```
SELECT * FROM INFORMATION_SCHEMA.INNODB_TABLESPACES;
```

-- *Check tablespace size*

```
SELECT  
    SPACE,  
    NAME,  
    FILE_TYPE,  
    INITIAL_SIZE,  
    CURRENT_SIZE  
FROM INFORMATION_SCHEMA.INNODB_TABLESPACES  
WHERE NAME = 'database_name/table_name';
```

3.5 Monitoring InnoDB with SHOW ENGINE INNODB STATUS

Running the Command

SHOW ENGINE INNODB STATUS\G

Output Sections

Sections in Output:

1. **Background Threads:** I/O threads, logger thread, master thread
2. **Semaphores:** Thread synchronization state
3. **Current Transaction State:** Active transactions and locks
4. **Pending I/O Reads/Writes:** Active I/O operations
5. **Buffer Pool and Memory:** Buffer pool usage statistics
6. **Transaction Log:** Log sequence number, checkpoint info

Key Metrics to Monitor

1. Buffer Pool Efficiency

Buffer pool hit rate = $(\text{buffer_pool_reads} / \text{buffer_pool_total_reads}) * 100$

Ideal: > 99%

Poor: < 98%

2. Transaction Activity

Number of active transactions

Transactions waiting for locks

3. Lock Conflicts

Number of lock waits

Deadlock detection and resolution

Interpreting InnoDB Status

2026-01-14 16:30:45 0x7f8b3c001700

Trx id counter 45

Purge done for trx's n:o < 42 undo n:o < 0

History list length 0

List of transactions for each session:

---TRANSACTION 42, ACTIVE 0 sec

2 lock struct(s), heap size 1136, 0 row lock(s), undo log entries 0

MySQL thread id 8, OS thread handle 0x7f8b3c0e5700, query id 123 localhost root

init

3.6 Summary: Key Takeaways

1. **ACID Properties:** Ensure data reliability and consistency
 - o Atomicity: All or nothing
 - o Consistency: Valid state to valid state
 - o Isolation: Transactions don't interfere
 - o Durability: Committed data persists
2. **Transaction Control:** BEGIN, COMMIT, ROLLBACK, SAVEPOINT
3. **Isolation Levels:** REPEATABLE READ (default), READ COMMITTED recommended
4. **Locking:** Row-level locks, gap locks, deadlock avoidance
5. **Tablespaces:** File-per-table recommended, monitor with INFORMATION_SCHEMA
6. **Monitoring:** Use SHOW ENGINE INNODB STATUS for health checks