

Advanced Replication & High Availability

8.1 InnoDB Cluster Architecture

What is InnoDB Cluster?

InnoDB Cluster is MySQL's built-in high-availability solution combining:

- **Group Replication:** Automatic replication with data consistency
- **MySQL Router:** Connection routing and load balancing
- **MySQL Shell:** Administration and automation

InnoDB Cluster Components

Group Replication:

- Distributed, fault-tolerant replication
- Multi-primary or single-primary mode
- Automatic failover on node failure
- Data consistency guaranteed

MySQL Router:

- Routes client connections to healthy nodes
- Handles failover automatically
- Read/write split for load balancing
- Session persistence

MySQL Shell:

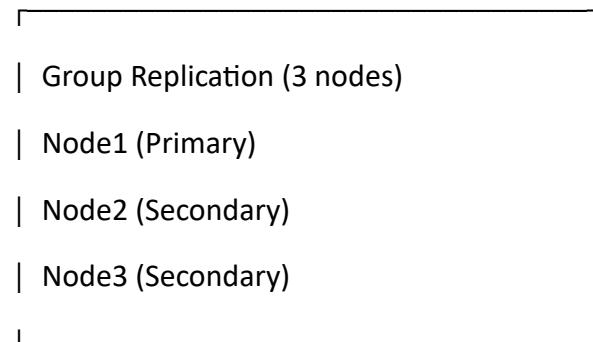
- Advanced command-line interface
- Cluster administration commands
- Admin API for automation
- JavaScript/Python scripting support

Cluster Topology

Application



MySQL Router (Connection Routing)



Data Data Data

8.2 Setting Up 3-Node InnoDB Cluster

Prerequisites

- Three MySQL 8.0.21+ servers
- Unique server_id for each
- InnoDB storage engine
- MySQL Shell installed
- Network connectivity between nodes

Initial Configuration

For each node (modify server_id):

[mysqld]

Unique server identifier

server_id = 1 # Change to 2, 3 for other nodes

Binary logging

```
log_bin = /var/log/mysql/mysql-bin  
binlog_format = ROW  
server_uuid = auto # Let MySQL generate  
  
# Group Replication  
transaction_write_set_extraction = XXHASH64  
loose_group_replication_group_name = "aaaaaaaa-bbbb-cccc-dddd-eeeeeeeeeee"  
loose_group_replication_start_on_boot = OFF  
loose_group_replication_local_address = "192.168.1.100:33061" # Unique per node  
loose_group_replication_group_seeds =  
"192.168.1.100:33061,192.168.1.101:33061,192.168.1.102:33061"  
loose_group_replication_single_primary_mode = ON # Single primary mode
```

InnoDB

```
default_table_encryption = OFF
```

Create InnoDB Cluster

Step 1: Connect to Primary Node

```
mysqlsh
```

In MySQL Shell:

```
// Connect to primary node  
shell> \connect icroot@192.168.1.100:3306
```

```
// Create cluster  
shell> var cluster = dba.createCluster("myCluster")  
  
// Cluster created successfully!  
// {
```

```

// "clusterName": "myCluster",
// "defaultReplicaSet": {
//   "name": "default",
//   "primary": "192.168.1.100:3306",
//   "status": "OK",
//   "statusText": "Cluster is ONLINE and can tolerate up to ONE failure.",
//   "topology": {
//     "192.168.1.100:3306": {
//       "address": "192.168.1.100:3306",
//       "label": "localhost",
//       "role": "HA",
//       "status": "ONLINE"
//     }
//   }
// }
// }
```

Step 2: Add Secondary Nodes

```

// Add second node
shell> cluster.addInstance('icroot@192.168.1.101:3306')
```

```

// Add third node
shell> cluster.addInstance('icroot@192.168.1.102:3306')
```

```

// Check cluster status
shell> cluster.status()
```

MySQL Router Configuration

```
# Bootstrap router (automatically configures)
sudo mysqlrouter --bootstrap icroot@192.168.1.100:3306 \
--directory /etc/mysqlrouter
```

Start MySQL Router

```
sudo systemctl start mysqlrouter
sudo systemctl enable mysqlrouter
```

Check status

```
sudo systemctl status mysqlrouter
```

Connect Through MySQL Router

Read-write port (6446) - connects to primary

```
mysql -h 127.0.0.1 -P 6446 -u icroot -p
```

Read-only port (6447) - connects to secondaries

```
mysql -h 127.0.0.1 -P 6447 -u icroot -p
```

X Protocol port (6448) - for advanced connections

```
mysql -h 127.0.0.1 -P 6448 -u icroot -p
```

8.3 Testing Automatic Failover

Create Test Database

```
CREATE DATABASE cluster_test;
USE cluster_test;
```

```
CREATE TABLE data (
    id INT PRIMARY KEY AUTO_INCREMENT,
    value VARCHAR(255),
    timestamp TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

```
INSERT INTO data (value) VALUES ('Initial data 1');
```

```
INSERT INTO data (value) VALUES ('Initial data 2');
```

```
SELECT * FROM data;
```

Simulate Primary Failure

Terminal 1: Monitor cluster

```
mysqlsh
```

```
shell> \connect icroot@192.168.1.100:3306
```

```
shell> var cluster = dba.getCluster()
```

```
shell> cluster.status()
```

Terminal 2: Stop primary node

On primary (192.168.1.100)

```
sudo systemctl stop mysql
```

Monitor automatic failover in Terminal 1

// Failover should occur automatically

```
shell> cluster.status()
```

// Should see:

// "primary": "192.168.1.101:3306" // New primary

```
// "status": "OK_NO_TOLERANCE" // Can't tolerate more failures
```

Verify Application Continuity

Applications should auto-failover to new primary

```
mysql -h 127.0.0.1 -P 6446 -u icroot -p \  
-e "SELECT * FROM cluster_test.data;"
```

New data should still be accessible

```
mysql -h 127.0.0.1 -P 6446 -u icroot -p \  
-e "INSERT INTO cluster_test.data (value) VALUES ('After failover');"
```

Restart Failed Node

Restart stopped node

```
sudo systemctl start mysql
```

In MySQL Shell, cluster should automatically detect

```
shell> cluster.status()
```

```
// Should show node coming back online
```

8.4 Replication Troubleshooting

Check Replication Status

-- On any node

```
SHOW SLAVE STATUS\G
```

```
SHOW REPLICA STATUS\G -- MySQL 8.0.22+ syntax
```

-- Key fields:
-- Slave_IO_Running: Yes
-- Slave_SQL_Running: Yes
-- Seconds_Behind_Master: 0

Monitor Group Replication

-- View group members
SELECT * FROM performance_schema.replication_group_members;

-- View group replication applier status
SELECT * FROM performance_schema.replication_applier_status\G

-- View group replication connection status
SELECT * FROM performance_schema.replication_connection_status\G

Common Replication Issues

Issue 1: Node Not Joining Cluster

-- Check error log
grep -i "group replication" /var/log/mysql/error.log

-- Common causes:
-- 1. Network connectivity
-- 2. Configuration mismatch
-- 3. Data conflicts

-- Solution: Check group_replication_group_seeds

Issue 2: High Replication Lag

-- Monitor applier queue

```
SELECT * FROM performance_schema.replication_applier_status_by_worker\G
```

-- Check for slow queries

```
SELECT * FROM performance_schema.events_statements_summary_by_digest
WHERE DIGEST_TEXT LIKE '%SELECT%'
ORDER BY SUM_TIMER_WAIT DESC
LIMIT 10;
```

Issue 3: Data Inconsistency

Use Percona Toolkit to check consistency

```
pt-table-checksum \
--host=192.168.1.100 \
--user=repl_user \
--password=ReplPass123!
```

Sync differences

```
pt-table-sync --execute \
h=192.168.1.100,u=repl_user,p=ReplPass123! \
h=192.168.1.101,u=repl_user,p=ReplPass123!
```

8.5 Cluster Administration

Add New Node to Running Cluster

mysqlsh

```
shell> \connect icroot@192.168.1.100:3306
```

```
shell> var cluster = dba.getCluster()
```

```
shell> cluster.addInstance('icroot@192.168.1.104:3306')
```

Remove Node from Cluster

```
shell> cluster.removeInstance('icroot@192.168.1.104:3306')
```

Change Primary Node

```
shell> cluster.setPrimaryInstance('icroot@192.168.1.102:3306')
```

// Waits for secondaries to catch up before switching

Upgrade Cluster

// Upgrade secondaries first, then primary

// Each node upgraded in rolling fashion

// No downtime

```
shell> cluster.upgradeMetadataAndGroupReplication()
```

Monitor Cluster Health

// Regular health check

```
shell> cluster.status()
```

// Detailed diagnostics

```
shell> cluster.describe()
```

// Check node status

```
shell> dba.getCluster().members
```

8.6 Summary: Key Takeaways

1. **InnoDB Cluster:** Three-component solution (Group Replication, MySQL Router, Shell)
2. **Automatic Failover:** Detects failures and promotes secondary automatically

3. **Setup:** 3 nodes with group replication, bootstrap router, connect via router
4. **Testing:** Simulate failures and verify automatic failover
5. **Troubleshooting:** Monitor group replication status, check data consistency
6. **Administration:** Add/remove nodes, change primary, rolling upgrades
7. **High Availability:** Reduces RTO and RPO for production databases