

# Hands On Lab: Advanced Replication & High Availability

## LAB 8.1: InnoDB Cluster Setup with 3 Nodes

**Objective:** Set up a production-grade InnoDB Cluster with automatic failover.

### Prerequisites:

- Three MySQL Server 8.0+ instances
- Node 1: 192.168.1.100:3306
- Node 2: 192.168.1.101:3306
- Node 3: 192.168.1.102:3306
- MySQL Shell installed

### Step-by-Step Instructions:

#### 1. Prepare MySQL Servers for InnoDB Cluster

For each server, configure as follows:

```
sudo vim /etc/mysql/mysql.conf.d/mysqld.cnf
```

Add under [mysqld]:

```
# InnoDB Cluster Configuration
```

```
server_id = 1 # Change to 2, 3 for other nodes
```

```
log_bin = /var/log/mysql/mysql-bin
```

```
binlog_format = ROW
```

```
# Group Replication
```

```
transaction_write_set_extraction = XXHASH64
```

```
loose_group_replication_group_name = "aaaaaaaa-bbbb-cccc-dddd-eeeeeeeeeee"
```

```
loose_group_replication_start_on_boot = OFF
```

```
loose_group_replication_local_address = "192.168.1.100:33061" # Change IP
```

```
loose_group_replication_group_seeds =
```

```
"192.168.1.100:33061,192.168.1.101:33061,192.168.1.102:33061"
```

```
loose_group_replication_single_primary_mode = ON
```

```
# InnoDB Cluster
```

```
default_table_encryption = OFF
```

## 2. Restart All MySQL Servers

```
sudo systemctl restart mysql
```

## 3. Create InnoDB Cluster Admin User on Primary

```
CREATE USER 'icroot'@'%' IDENTIFIED BY 'IcPass123!';  
GRANT ALL ON *.* TO 'icroot'@'%' WITH GRANT OPTION;  
CREATE USER 'icroot'@'localhost' IDENTIFIED BY 'IcPass123!';  
GRANT ALL ON *.* TO 'icroot'@'localhost' WITH GRANT OPTION;  
FLUSH PRIVILEGES;
```

## 4. Initialize InnoDB Cluster Using MySQL Shell

```
mysqlsh
```

In MySQL Shell:

```
// Connect to primary node  
shell> \connect icroot@192.168.1.100:3306
```

```
// Create cluster  
shell> var cluster = dba.createCluster("myCluster")
```

```
// Add members to cluster  
shell> cluster.addInstance('icroot@192.168.1.101:3306')  
shell> cluster.addInstance('icroot@192.168.1.102:3306')
```

```
// Get cluster status
```

```
shell> cluster.status()
```

```
// Exit MySQL Shell
```

```
shell> \quit
```

## 5. Deploy MySQL Router for Connection Pooling

```
# Install MySQL Router
```

```
sudo apt install -y mysql-router-community
```

```
# Configure router
```

```
sudo mysqlrouter --bootstrap icroot@192.168.1.100:3306 \
```

```
--directory /etc/mysqlrouter
```

```
# Start router
```

```
sudo systemctl start mysqlrouter
```

```
sudo systemctl status mysqlrouter
```

## 6. Connect Through MySQL Router

```
# Connect to read-write port (6446)
```

```
mysql -h 127.0.0.1 -P 6446 -u icroot -p
```

```
# Connect to read-only port (6447)
```

```
mysql -h 127.0.0.1 -P 6447 -u icroot -p
```

## 7. Test Automatic Failover

### a) Create test data:

```
CREATE DATABASE cluster_test;
```

```
CREATE TABLE cluster_test.data (
```

```
    id INT PRIMARY KEY AUTO_INCREMENT,
```

```
    value VARCHAR(255)
```

```
) ENGINE=InnoDB;
```

```
INSERT INTO cluster_test.data (value) VALUES ('Test data 1'), ('Test data 2');
```

**b) Simulate primary failure:**

*# On primary node, stop MySQL*

```
sudo systemctl stop mysql
```

**c) Verify automatic failover:**

*# In MySQL Shell, check cluster status*

```
mysqlsh
```

```
shell> \connect icroot@192.168.1.101:3306
```

```
shell> var cluster = dba.getCluster()
```

```
shell> cluster.status()
```

**d) Verify data availability:**

*# Connect through router*

```
mysql -h 127.0.0.1 -P 6446 -u icroot -p -e "SELECT * FROM cluster_test.data;"
```

## 8. Monitor Cluster Health

```
mysqlsh
```

```
shell> \connect icroot@192.168.1.101:3306
```

```
shell> var cluster = dba.getCluster()
```

```
shell> cluster.describe()
```

```
shell> cluster.status()
```

### Hands-on Tasks:

- Set up 3-node InnoDB Cluster
- Configure MySQL Router
- Test automatic failover
- Monitor cluster status
- Document failover procedures

**Solution:**

1. Set up a 3-node MySQL InnoDB Cluster using MySQL Shell AdminAPI with one primary and two secondary instances.

**1. Prerequisites and instance layout**

Use three MySQL 8.x instances (ports are examples):

- Node1: localhost:3310 (seed / primary)
- Node2: localhost:3320
- Node3: localhost:3330

**Ensure on all three:**

```
[mysqld]
server_id=1/2/3      # unique per node
gtid_mode=ON
enforce_gtid_consistency=ON
binlog_format=ROW
log_bin=binlog
transaction_write_set_extraction=XXHASH64
loose-group_replication_bootstrap_group=OFF
loose-group_replication_start_on_boot=OFF
loose-group_replication_group_name="aaaaaaaa-bbbb-cccc-dddd-eeeeeeeeeee"
loose-group_replication_local_address="127.0.0.1:33061/33062/33063"
loose-group_replication_group_seeds="127.0.0.1:33061,127.0.0.1:33062,127.0.0.1:33063"
loose-group_replication_single_primary_mode=ON
loose-group_replication_enforce_update_everywhere_checks=OFF
```

Restart all instances after editing config.

**2. Create cluster admin user**

On each node (or on Node1 and let AdminAPI create internal accounts):

```
CREATE USER 'cluster_admin'@'%' IDENTIFIED BY 'ClusterPass2026!';
```

```
GRANT ALL PRIVILEGES ON *.* TO 'cluster_admin'@'%' WITH GRANT OPTION;  
FLUSH PRIVILEGES;
```

### **3. Check and configure instances with MySQL Shell**

**Start MySQL Shell in JS mode:**

```
mysqlsh --js
```

**On Node1:**

```
shell.connect('cluster_admin@localhost:3310');
```

```
dba.checkInstanceConfiguration('cluster_admin@localhost:3310');
```

```
dba.configureInstance('cluster_admin@localhost:3310');
```

```
dba.checkInstanceConfiguration('cluster_admin@localhost:3320');
```

```
dba.configureInstance('cluster_admin@localhost:3320');
```

```
dba.checkInstanceConfiguration('cluster_admin@localhost:3330');
```

```
dba.configureInstance('cluster_admin@localhost:3330');
```

Confirm checks report “instance configuration is suitable” or that required changes were applied.

### **4. Create the InnoDB Cluster**

**Still in MySQL Shell connected to Node1:**

```
var cluster = dba.createCluster(  
    'TrainingCluster',  
    { localAddress: 'localhost:3310' }  
)
```

This:

- Deploys cluster metadata on Node1
- Configures Group Replication and starts it on Node1

Check status:

```
cluster.status();
```

Node1 should show as PRIMARY with mode R/W.

## 5. Add the other two nodes

**From the same shell:**

```
cluster.addInstance('cluster_admin@localhost:3320', {  
  localAddress: 'localhost:3320',  
  recoveryMethod: 'clone'  
});
```

```
cluster.addInstance('cluster_admin@localhost:3330', {  
  localAddress: 'localhost:3330',  
  recoveryMethod: 'clone'  
});
```

**Wait for each clone to finish. Then:**

```
cluster.status({extended: true});
```

You should see:

- Node1: PRIMARY, ONLINE
- Node2: SECONDARY, ONLINE
- Node3: SECONDARY, ONLINE

## 6. Basic functional tests

### 6.1 Test writes on primary and reads on all nodes

From MySQL Shell, connect to primary:

```
shell.connect('cluster_admin@localhost:3310');

session.runSql("CREATE DATABASE IF NOT EXISTS cluster_lab;");

session.runSql("USE cluster_lab;");

session.runSql(`

CREATE TABLE employees (

    id INT PRIMARY KEY AUTO_INCREMENT,
    name VARCHAR(100),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
)

`);

session.runSql("INSERT INTO employees (name) VALUES ('Node1-insert');");
```

#### On Node2 and Node3 (SQL mode):

```
shell.connect('cluster_admin@localhost:3320');

session.runSql("SELECT @@server_id, @@read_only, COUNT(*) FROM
cluster_lab.employees;");
```

```
shell.connect('cluster_admin@localhost:3330');

session.runSql("SELECT @@server_id, @@read_only, COUNT(*) FROM
cluster_lab.employees;");
```

Row counts should match and secondaries should be read\_only = 1.

### 6.2 Test automatic primary election

#### Simulate failure of Node1:

```
# Stop mysqld on Node1
```

```
sudo systemctl stop mysql@3310
```

**Back in MySQL Shell (connected to one of surviving nodes):**

```
shell.connect('cluster_admin@localhost:3320');
cluster = dba.getCluster('TrainingCluster');
cluster.status();
```

One of Node2/Node3 should now be PRIMARY and writable.

## 7. (Optional) Sandbox deployment helper

**For purely local hands-on, AdminAPI can create sandboxes for you:**

```
mysqlsh --dba deploy-sandbox-instance 3310 --clusterAdmin=cluster_admin --
clusterAdminPassword=ClusterPass2026!
```

```
mysqlsh --dba deploy-sandbox-instance 3320 --clusterAdmin=cluster_admin --
clusterAdminPassword=ClusterPass2026!
```

```
mysqlsh --dba deploy-sandbox-instance 3330 --clusterAdmin=cluster_admin --
clusterAdminPassword=ClusterPass2026!
```

**Then:**

```
shell.connect('cluster_admin@localhost:3310');
var cluster = dba.createCluster('TrainingCluster');
cluster.addInstance('cluster_admin@localhost:3320');
cluster.addInstance('cluster_admin@localhost:3330');
cluster.status();
```

## 8. Verification checklist

- All 3 instances have GTID, binlog, and group\_replication settings correct.
- dba.checkInstanceConfiguration() passes for each node.
- cluster.status() shows 1 PRIMARY, 2 SECONDARY in ONLINE state.
- Writes on primary replicate to both secondaries.
- Stopping primary triggers automatic primary election on a secondary.

This completes a full 3-node InnoDB Cluster lab setup using MySQL Shell AdminAPI.

**2. Configure MySQL Router to front your 3-node InnoDB Cluster and expose stable R/W and R/O endpoints for applications.**

### **1. Install MySQL Router**

**On the application (or separate) host:**

```
# Ubuntu / Debian
```

```
sudo apt update
```

```
sudo apt install mysql-router
```

```
# Verify
```

```
mysqlrouter --version
```

### **2. Create router account in the cluster**

**In MySQL Shell connected to the cluster primary (e.g., localhost:3310):**

```
mysqlsh --js
```

```
shell.connect('cluster_admin@localhost:3310');
```

```
cluster = dba.getCluster('TrainingCluster');
```

```
// Create a low-privilege router account
```

```
cluster.setupRouterAccount('router_user', {password: 'RouterPass2026!'});
```

This creates the account and grants only the privileges needed for metadata and health checks.

### **3. Bootstrap MySQL Router**

**On the router host (can be same as primary for a lab):**

```
cd ~ # or /opt  
mysqlrouter \  
--bootstrap router_user@localhost:3310 \  
--directory myrouter \  
--user $USER
```

You will be prompted for RouterPass2026!.

**Bootstrap will:**

- Connect to the InnoDB Cluster via the specified node.
- Read metadata and topology.
- Create myrouter/mysqlrouter.conf.
- Generate helper scripts start.sh, stop.sh (in sandbox mode).

### **4. Inspect generated configuration**

```
cd myrouter  
ls  
# mysqlrouter.conf, start.sh, stop.sh, run/ logs/ etc.
```

**Open mysqlrouter.conf:**

```
[DEFAULT]  
  
user = youruser  
  
logging_folder = /home/youruser/myrouter/log  
  
[logger]  
  
level = INFO  
  
[routing:primary_rw]
```

```
bind_address = 0.0.0.0
bind_port = 6446
destinations = metadata-cache://TrainingCluster/?role=PRIMARY
routing_strategy = first-available
```

```
[routing:secondary_ro]
bind_address = 0.0.0.0
bind_port = 6447
destinations = metadata-cache://TrainingCluster/?role=SECONDARY
routing_strategy = round-robin
```

#### **Ports (default sandbox):**

- 6446 – Classic protocol R/W
- 6447 – Classic protocol R/O
- 6448 – X protocol R/W
- 6449 – X protocol R/O

## **5. Start MySQL Router**

#### **For lab/sandbox mode:**

```
cd ~/myrouter
./start.sh # or: mysqlrouter -c mysqlrouter.conf &
```

#### **Check it is listening:**

```
ss -ltnp | grep 644
# or
netstat -ltnp | grep 644
```

## 6. Test routing behavior

### 6.1 Read-write endpoint

# Connect through router R/W port

```
mysql -h 127.0.0.1 -P 6446 -u app_user -p
```

```
SELECT @@hostname, @@port, @@global.read_only;
```

-- Should show current PRIMARY node, read\_only = 0

```
USE cluster_lab;
```

```
INSERT INTO employees(name) VALUES('via_router_rw');
```

```
SELECT COUNT(*) FROM employees;
```

### 6.2 Read-only endpoint

```
mysql -h 127.0.0.1 -P 6447 -u app_user -p
```

```
SELECT @@hostname, @@port, @@global.read_only;
```

-- Should show a SECONDARY, read\_only = 1

```
SELECT * FROM cluster_lab.employees ORDER BY id DESC LIMIT 5;
```

-- Read succeeds; writes fail:

```
INSERT INTO employees(name) VALUES('should_fail');
```

-- ERROR: read-only or access denied

## 7. Validate failover with Router

### 1. Confirm current primary via router:

```
mysql -h 127.0.0.1 -P 6446 -u app_user -p -e "SELECT @@hostname, @@port, @@global.read_only;"
```

**2. Stop the current primary MySQL instance (e.g., 3310):**

```
sudo systemctl stop mysql@3310
```

**3. Wait for Group Replication to elect a new primary, then test again:**

```
mysql -h 127.0.0.1 -P 6446 -u app_user -p -e "SELECT @@hostname, @@port,  
@@global.read_only;"
```

You should see a different host/port, with read\_only = 0. Router transparently redirects R/W traffic to the new primary.

**8. Optional: Router as a daemon/service**

**To run Router as a service using the generated config:**

```
sudo mysqlrouter -c /home/youruser/myrouter/mysqlrouter.conf --user mysqlrouter_user --  
daemon
```

Or create a systemd unit pointing to mysqlrouter -c /path/to/mysqlrouter.conf.

**9. Lab verification checklist**

- cluster.setupRouterAccount() executed successfully.
- mysqlrouter --bootstrap completed and mysqlrouter.conf created.
- R/W port 6446 routes to primary, writes succeed.
- R/O port 6447 routes to secondaries, writes blocked.
- After primary shutdown, Router redirects R/W connections to new primary.
- Logs (myrouter/log/mysqlrouter.log) show healthy metadata refresh and no fatal errors.

This completes a full MySQL Router configuration lab for your 3-node InnoDB Cluster.

**3.** Test automatic failover in your 3-node InnoDB Cluster and verify that MySQL Router transparently reconnects clients to the new primary.

**1. Pre-check: cluster and router health**

**In MySQL Shell (JS), connect via Router R/W port:**

```
mysqlsh --js root@localhost:6446
```

**Check cluster status:**

```
cluster = dba.getCluster('TrainingCluster');  
cluster.status({extended: true});
```

**Confirm:**

- 3 members ONLINE
- One PRIMARY, two SECONDARY
- singlePrimaryMode: true

Note the current primary, e.g. localhost:3310.

**2. Open a client session through Router**

**From a separate terminal, connect through the R/W port:**

```
mysql -h 127.0.0.1 -P 6446 -u app_user -p
```

**In that session:**

```
SELECT @@hostname AS server, @@port AS port, @@global.read_only AS read_only;
```

*-- Should show current PRIMARY with read\_only = 0*

Keep this client session open; you will use it to observe failover.

### **3. Start a continuous workload**

**Still on the client through Router:**

```
USE cluster_lab;
```

-- Create simple table if not present

```
CREATE TABLE IF NOT EXISTS ha_test (
    id INT PRIMARY KEY AUTO_INCREMENT,
    msg VARCHAR(100),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
) ENGINE=InnoDB;
```

-- Simple loop (in bash) to generate writes via Router

```
\! while true; do \
    mysql -h 127.0.0.1 -P 6446 -u app_user -p'yourpass' -D cluster_lab \
    -e "INSERT INTO ha_test(msg) VALUES(CONCAT('write-', NOW()));" \
    sleep 1; \
done
```

This simulates a live application writing every second.

### **4. Trigger automatic failover by stopping the primary**

**From another terminal on the primary host (e.g. instance on 3310):**

```
# Stop only the primary MySQL instance
sudo systemctl stop mysql@3310 # or appropriate service name
```

**Alternatively (lab, not prod):**

```
mysqladmin -h 127.0.0.1 -P 3310 -u root -p shutdown
```

Group Replication detects the loss and automatically elects a new primary from remaining members.

## 5. Observe behavior from the client

### In the Router-connected client:

- You may see one transient error (connection lost or timeout) during election.
- Then inserts resume automatically once Router reconnects to the new primary.

### Confirm which server you are now writing to:

```
SELECT @@hostname AS server, @@port AS port, @@global.read_only AS read_only;
```

```
-- Should now show one of 3320/3330 with read_only = 0
```

```
SELECT COUNT(*) FROM ha_test;
```

The count should continue increasing; there must be no manual change in connection string (still pointing to 127.0.0.1:6446).

## 6. Verify cluster sees the new primary

### Back in MySQL Shell:

```
cluster = dba.getCluster('TrainingCluster');  
cluster.status({extended: true});
```

### You should see:

- Old primary (3310) as UNREACHABLE or OFFLINE
- A surviving member (3320 or 3330) promoted to PRIMARY
- Remaining node as SECONDARY

## 7. Optional: Bring back the old primary

### Restart the stopped instance:

```
sudo systemctl start mysql@3310
```

### Then rejoin to the cluster (if not automatic):

```
shell.connect('cluster_admin@localhost:3320'); // current primary
```

```
cluster = dba.getCluster('TrainingCluster');

cluster.rejoinInstance('cluster_admin@localhost:3310');

cluster.status();
```

Old primary should come back as SECONDARY.

## 8. Failover test checklist

- Before test: 3 nodes ONLINE, 1 PRIMARY, 2 SECONDARY.
- Router R/W port 6446 returns PRIMARY and accepts writes.
- After stopping primary, cluster.status() shows new PRIMARY elected.
- Router-connected client resumes writes without changing connection string.
- SELECT @@hostname via Router shows new PRIMARY.
- Bringing back old primary re-adds it as SECONDARY, not PRIMARY.

This lab confirms automatic failover of InnoDB Cluster and transparent reconnection through MySQL Router.

4. Monitor MySQL InnoDB Cluster status using MySQL Shell's AdminAPI, including real-time watch and key health indicators.

### 1. Basic cluster status check

Connect MySQL Shell (JS) to any ONLINE instance (preferably the primary):

```
mysqlsh --js cluster_admin@localhost:3310
```

Get a cluster handle and show status:

```
var cluster = dba.getCluster('TrainingCluster'); // or dba.getCluster()  
cluster.status();
```

Typical output (trimmed):

```
{  
  "clusterName": "TrainingCluster",  
  "defaultReplicaSet": {  
    "primary": "localhost:3310",  
    "status": "OK",  
    "statusText": "Cluster is ONLINE and can tolerate up to ONE failure.",  
    "topology": {  
      "localhost:3310": { "mode": "R/W", "status": "ONLINE" },  
      "localhost:3320": { "mode": "R/O", "status": "ONLINE" },  
      "localhost:3330": { "mode": "R/O", "status": "ONLINE" }  
    }  
  }  
}
```

Key fields:

- status: OK, OK\_NO\_TOLERANCE, DEGRADED, UNREACHABLE, etc.

- `statusText`: human explanation (tolerance, failures).
- `primary`: current primary (single-primary mode).
- `topology`: each member's status (ONLINE, OFFLINE, RECOVERING) and mode (R/W, R/O).

## 2. Extended status for deeper monitoring

**Use extended mode for more detail:**

```
cluster.status({extended: 1});
```

**Higher levels (1–3) add:**

- recovery information (clone or incremental, progress %).
- `applierStatus` and `receiverStatus` for Group Replication threads.
- GTID and replication lag indicators per instance.

**For full detail:**

```
cluster.status({extended: 3});
```

Use this during node recovery or troubleshooting.

## 3. Real-time monitoring with \watch

**MySQL Shell can auto-refresh any command output:**

```
\watch 5 cluster.status()
```

- Refreshes cluster status every 5 seconds.
- Use when simulating failures (stopping a node, testing failover) to observe topology changes live.

**Example during a node outage:**

- `statusText` changes to "Cluster is ONLINE but can only tolerate up to ZERO failures."
- One member shows status: "UNREACHABLE" or OFFLINE.

## 4. Quick CLI health checks

**For scripting from Bash (no interactive JS):**

```
mysqlsh --js cluster_admin@localhost:3310 -e "dba.getCluster('TrainingCluster').status();" > cluster_status.json
```

**Or a concise health summary:**

```
mysqlsh --js cluster_admin@localhost:3310 -e "
var c = dba.getCluster('TrainingCluster');
var s = c.status();
print('Status: ' + s.defaultReplicaSet.status);
print('Primary: ' + s.defaultReplicaSet.primary);
"
```

Use this in cron or monitoring agents.

## 5. What to monitor (lab checklist)

**Run regularly:**

```
var s = cluster.status();
s.defaultReplicaSet.status;      // Should be "OK"
s.defaultReplicaSet.primary;    // Current primary
s.defaultReplicaSet.topology;   // Per-node status/mode
```

**In your lab, verify:**

- status is OK, statusText says cluster is ONLINE and can tolerate at least one failure.
- Exactly one node shows mode: "R/W" (primary), others R/O.
- When you stop one node, status reflects reduced tolerance and node status changes.
- When node is restarted and rejoined, all return to ONLINE.

This gives a complete, hands-on monitoring solution for your InnoDB Cluster using MySQL Shell.

## **LAB 8.2: Replication Monitoring and Troubleshooting**

**Objective:** Monitor replication and troubleshoot common issues.

### **Step-by-Step Instructions:**

#### **1. Monitor Replication Status**

-- *On Slave*

```
SHOW SLAVE STATUS\G
```

-- *Key fields to monitor:*

-- *Slave\_IO\_Running: Yes*

-- *Slave\_SQL\_Running: Yes*

-- *Seconds\_Behind\_Master: 0 (or acceptable lag)*

-- *Last\_Errno: 0 (no errors)*

-- *Last\_Error: (should be empty)*

#### **2. Check Master Binary Log Position**

-- *On Master*

```
SHOW MASTER STATUS;
```

```
SHOW BINARY LOGS;
```

#### **3. Monitor Slave Replication Threads**

-- *On Slave*

```
SHOW PROCESSLIST;
```

-- *Look for:*

-- *Slave I/O: Reading event from the master*

-- *Slave SQL: Executing event received from master*

#### **4. Check Replication Lag**

-- *Method 1: Using SHOW SLAVE STATUS*

```
SHOW SLAVE STATUS\G
```

-- *Look for: Seconds\_Behind\_Master*

```
-- Method 2: Using Performance Schema  
SELECT * FROM performance_schema.replication_applier_status\G
```

## 5. Simulate Replication Error

### a) On Slave, insert conflicting data:

```
INSERT INTO replication_test.test_table (id, data)  
VALUES (1, 'Conflicting data on slave');
```

### b) On Master, update same row:

```
UPDATE replication_test.test_table SET data = 'Updated on master'  
WHERE id = 1;
```

### c) Check for errors on Slave:

```
SHOW SLAVE STATUS\G
```

-- Look for error in Last\_Error field

## 6. Skip Replication Error (if needed)

-- Option 1: Skip specific number of events

```
SET GLOBAL SQL_SLAVE_SKIP_COUNTER = 1;  
START SLAVE;
```

-- Option 2: Using GTID (preferred)

-- Don't use SQL\_SLAVE\_SKIP\_COUNTER with GTID

-- Instead, use:

```
SET GTID_NEXT='current_master_gtid:next_number';  
BEGIN;  
COMMIT;  
SET GTID_NEXT='AUTOMATIC';
```

## 7. Use Percona Toolkit for Replication Monitoring

# Install Percona Toolkit

```
sudo apt install -y percona-toolkit
```

```
# Check table consistency

pt-table-checksum --host=192.168.1.100 \
    --user=root --password=YourPassword \
    --databases=replication_test \
    --replicate=percona.checksums

# View results

mysql -u root -p -e "SELECT * FROM percona.checksums;"
```

```
# Fix differences

pt-table-sync --execute \
    h=192.168.1.100,u=root,p=YourPassword \
    --replicate=percona.checksums
```

## 8. Monitor with Performance Schema

-- View replication channels

```
SELECT * FROM performance_schema.replication_connection_configuration\G
SELECT * FROM performance_schema.replication_applier_configuration\G
```

-- View connection status

```
SELECT * FROM performance_schema.replication_connection_status\G
SELECT * FROM performance_schema.replication_applier_status\G
```

-- View worker status (for parallel replication)

```
SELECT * FROM performance_schema.replication_applier_status_by_worker\G
```

## 9. Set up Automated Monitoring

# Create monitoring script

```
cat > /usr/local/bin/monitor_replication.sh << 'EOF'
```

```

#!/bin/bash

SLAVE_HOST="192.168.1.101"
SLAVE_USER="root"
SLAVE_PASS="YourPassword"

SECONDS_BEHIND=$(mysql -h$SLAVE_HOST -u$SLAVE_USER -p$SLAVE_PASS \
-sNe "SHOW SLAVE STATUS\G" | grep "Seconds_Behind_Master" | awk '{print $NF}')

IO_RUNNING=$(mysql -h$SLAVE_HOST -u$SLAVE_USER -p$SLAVE_PASS \
-sNe "SHOW SLAVE STATUS\G" | grep "Slave_IO_Running" | awk '{print $NF}')

SQL_RUNNING=$(mysql -h$SLAVE_HOST -u$SLAVE_USER -p$SLAVE_PASS \
-sNe "SHOW SLAVE STATUS\G" | grep "Slave_SQL_Running" | awk '{print $NF}')

echo "$(date): IO=$IO_RUNNING, SQL=$SQL_RUNNING, Lag=${SECONDS_BEHIND}s"

if [[ "$IO_RUNNING" != "Yes" ]] || [[ "$SQL_RUNNING" != "Yes" ]]; then
    echo "ERROR: Replication issue detected!" | mail -s "Replication Alert"
    admin@example.com
fi
EOF

chmod +x /usr/local/bin/monitor_replication.sh

# Add to crontab
crontab -e

```

```
# Add: */5 * * * * /usr/local/bin/monitor_replication.sh >>
/var/log/replication_monitor.log
```

### Hands-on Tasks:

- Monitor replication continuously
- Simulate and resolve replication errors
- Check data consistency between master and slave
- Document monitoring procedures

### Solution:

1. Monitor MySQL replication continuously using a combination of SHOW REPLICAS STATUS, Performance Schema, and a lightweight watch/alert script.

#### 1. Simple continuous monitor using MySQL Shell

On the replica:

```
mysqlsh --sql root@replica_host:3306
```

Then:

```
\watch 5 SHOW REPLICAS STATUS\G
```

This:

- Refreshes replication status every 5 seconds.
- Lets you watch Replica\_IO\_Running, Replica\_SQL\_Running, Seconds\_Behind\_Source live.

Key fields to keep an eye on:

- Replica\_IO\_Running / Slave\_IO\_Running = Yes
- Replica\_SQL\_Running / Slave\_SQL\_Running = Yes
- Seconds\_Behind\_Source / Seconds\_Behind\_Master ≈ 0

## 2. Performance Schema-based monitoring query

For more detailed, queryable metrics (MySQL 8+):

**SELECT**

```
rcs.channel_name,  
rcs.service_state AS io_state,  
rcs.last_heartbeat_timestamp,  
ras.service_state AS sql_state,  
ras.last_error_number,  
ras.last_error_message
```

**FROM** performance\_schema.replication\_connection\_status rcs

**JOIN** performance\_schema.replication\_applier\_status ras

**ON** rcs.channel\_name = ras.channel\_name;

- Run this periodically (every N seconds) to track whether I/O and SQL threads are running and see last error messages.

## 3. Lightweight Bash monitor loop

Create monitor\_replication.sh on the replica:

```
#!/bin/bash  
  
HOST="127.0.0.1"  
  
USER="root"  
  
PASS="yourpass"  
  
INTERVAL=5  
  
MAX_LAG=30 # seconds
```

**while** true; **do**

```
STATUS=$(mysql -h"$HOST" -u"$USER" -p"$PASS" -e "SHOW REPLICAS STATUS\G"  
2>/dev/null)
```

```
IO=$(echo "$STATUS" | grep -E 'Replica_IO_Running:|Slave_IO_Running:' | awk '{print $2}')
```

```

SQL=$(echo "$STATUS" | grep -E 'Replica_SQL_Running:|Slave_SQL_Running:' | awk '{print $2}')

LAG=$(echo "$STATUS" | grep -E 'Seconds_Behind_Source:|Seconds_Behind_Master:' |
awk '{print $2}')


TS=$(date '+%Y-%m-%d %H:%M:%S')

echo "[${TS}] IO=$IO SQL=$SQL LAG=${LAG}s"

if [ "$IO" != "Yes" ] || [ "$SQL" != "Yes" ]; then

    echo "[${TS}] CRITICAL: Replication threads stopped! IO=$IO SQL=$SQL" | logger -t mysql-repl

    elif [ "$LAG" = "NULL" ] || [ "$LAG" -gt "$MAX_LAG" ]; then

        echo "[${TS}] WARNING: Replication lag high or unknown (LAG=$LAG)" | logger -t mysql-repl

    fi

sleep "$INTERVAL"

done

chmod +x monitor_replication.sh

./monitor_replication.sh

```

This gives a continuous stream of replication health lines with basic alerting via syslog.

#### **4. Lag-focused continuous monitor (Performance Schema)**

For multi-threaded replicas, use Performance Schema for more accurate lag:

```
\watch 5
```

```
SELECT
```

```
channel_name,  
COUNT_TRANSACTIONS_REMOTE_IN_APPLIER_QUEUE AS tx_in_queue  
FROM performance_schema.replication_applier_status_by_coordinator;
```

- If tx\_in\_queue grows and stays high, your replica is falling behind even if Seconds\_Behind\_Source looks small.

#### **5. Lab verification steps**

In your lab, do this sequence:

1. Start monitor\_replication.sh on the replica.
2. On the master, run a write loop:

```
while true; do
```

```
    mysql -u root -p -e "INSERT INTO test_backup_restore.rep_monitor(msg) VALUES(NOW());"  
    sleep 1
```

```
done
```

3. Watch the monitor output:

- IO(SQL Yes, lag stays at 0–1 seconds).

4. Temporarily stop SQL thread on replica:

```
STOP REPLICA SQL_THREAD;
```

5. Monitor should show:

- SQL=No, lag becomes NULL, critical alert lines.

6. Start SQL thread again:

```
START REPLICA SQL_THREAD;
```

7. Monitor returns to healthy state.

That gives you a complete, hands-on continuous replication monitoring setup suitable for lab and easily adaptable to production.

2. Simulate common MySQL replication errors (duplicate key and missing row) and practice safe ways to diagnose and fix them without corrupting data.

### 1. Setup: simple replicated table

On master:

```
CREATE DATABASE IF NOT EXISTS repl_lab;
```

```
USE repl_lab;
```

```
CREATE TABLE accounts (
    id INT PRIMARY KEY,
    balance INT NOT NULL
) ENGINE=InnoDB;
```

```
INSERT INTO accounts VALUES (1,100), (2,200), (3,300);
```

Verify data on slave:

```
USE repl_lab;
```

```
SELECT * FROM accounts ORDER BY id;
```

## 2. Simulate error 1062: duplicate entry on slave

### 2.1 Manually create conflicting row on slave

On slave (bypass replication, simulate bad manual write):

```
USE repl_lab;
```

```
INSERT INTO accounts VALUES (4,400); -- Create row that master does NOT have
```

### 2.2 Generate conflicting event from master

On master:

```
USE repl_lab;
```

```
INSERT INTO accounts VALUES (4,500); -- Replicates to slave and conflicts
```

On slave, check status:

```
SHOW REPLICAS\G -- or SHOW SLAVE STATUS\G
```

You should see:

- Last\_SQL\_Errno: 1062
- Last\_SQL\_Error: Duplicate entry '4' for key 'PRIMARY' ...
- Replica\_SQL\_Running: No

## 3. Fix error 1062 safely (preferred)

Goal: make slave match master, then resume replication.

### 3.1 Inspect offending row and relay log position

On slave:

```
SHOW REPLICAS\G
```

-- Note: Relay\_Master\_Log\_File, Exec\_Master\_Log\_Pos, Last\_SQL\_Error

```
SELECT * FROM repl_lab.accounts WHERE id = 4;
```

You now know slave has a wrong id=4 row.

### 3.2 Remove bad row and restart SQL thread

On slave:

```
STOP REPLICA SQL_THREAD; -- or STOP SLAVE SQL_THREAD;
```

```
DELETE FROM repl_lab.accounts WHERE id = 4;
```

```
START REPLICA SQL_THREAD; -- or START SLAVE SQL_THREAD;
```

Check status again:

```
SHOW REPLICA STATUS\G
```

-- Last\_SQL\_Error: 0

-- Replica\_SQL\_Running: Yes

Verify:

```
SELECT * FROM repl_lab.accounts ORDER BY id;
```

-- Slave now has id=4 with balance=500, matching master

This is safer than skipping events because you fix the data and let the event replay correctly.

### 4. Simulate error 1032: “row not found” on slave

Now simulate the reverse: slave is missing a row master thinks exists.

#### 4.1 Create row on master

On master:

```
USE repl_lab;
```

```
INSERT INTO accounts VALUES (5,500);
```

Wait for replication, then on slave:

```
SELECT * FROM accounts WHERE id=5; -- Confirm present
```

## **4.2 Delete row only on slave**

On slave:

```
DELETE FROM accounts WHERE id=5;
```

## **4.3 Update row on master**

On master:

```
UPDATE accounts SET balance = balance + 50 WHERE id=5;
```

Replication tries to update id=5 on slave, which no longer exists.

On slave:

```
SHOW REPLICA STATUS\G
```

You should see:

- Last\_SQL\_Error: 1032
- Last\_SQL\_Error: Could not execute Update\_rows event ... Can't find record in 'accounts'
- Replica\_SQL\_Running: No

## **5. Fix error 1032 safely**

### **5.1 Capture failing event information**

On slave:

```
SHOW REPLICA STATUS\G
```

-- Note: Last\_SQL\_Error, Relay\_Master\_Log\_File, Exec\_Master\_Log\_Pos

### **5.2 Recreate missing row to match master**

Check master for correct data:

-- On master

```
SELECT * FROM repl_lab.accounts WHERE id=5;
```

-- Suppose result is (5,550)

On slave:

```
STOP REPLICA SQL_THREAD;
```

```
INSERT INTO repl_lab.accounts VALUES (5,550);
```

```
START REPLICA SQL_THREAD;
```

Verify:

```
SHOW REPLICA STATUS\G
```

```
-- Last_SQL_Errno: 0, Replica_SQL_Running: Yes
```

```
SELECT * FROM repl_lab.accounts WHERE id=5;
```

## 6. When (and when NOT) to use sql\_slave\_skip\_counter

You can also skip a single failing event, but this risks divergence if misused.

On slave:

```
STOP REPLICA SQL_THREAD;
```

```
SET GLOBAL sql_slave_skip_counter = 1;
```

```
START REPLICA SQL_THREAD;
```

Use this only when:

- You have confirmed the event is safe to ignore (e.g., duplicate insert that's already applied and data matches master).
- Ideally after validating the table with a tool like pt-table-checksum.

Do not use it blindly in loops or with slave-skip-errors in my.cnf in production.

## 7. Lab summary script

You can turn the above into a repeatable lab script for students:

1. Prepare table and initial data.

2. Run a shell script that:

- Performs conflicting manual changes on slave.
- Applies changes on master.
- Shows SHOW REPLICAS STATUS\G.
- Guides them through:
  - STOP REPLICA SQL\_THREAD
  - Fix data (DELETE / INSERT)
  - START REPLICA SQL\_THREAD

3. End with a consistency check:

-- On both master and slave

```
SELECT COUNT(*) AS rows, SUM(balance) AS total_balance FROM repl_lab.accounts;
```

If counts and totals match, replication is healthy again.

This lab walks through both simulating and safely resolving the two most common replication errors: 1062 and 1032, without resorting to dangerous global skip settings.

3. Check data consistency between master and slave by comparing counts and checksums at SQL level, then validating with pt-table-checksum for full-table drift detection.

## 1. Quick SQL-level consistency checks

### 1.1 Row counts per table

On master:

```
USE repl_lab; -- or your application DB
```

```
SELECT 'master' AS server, table_name, table_rows  
FROM information_schema.tables  
WHERE table_schema = 'repl_lab'  
ORDER BY table_name;
```

On slave:

```
USE repl_lab;
```

```
SELECT 'slave' AS server, table_name, table_rows  
FROM information_schema.tables  
WHERE table_schema = 'repl_lab'  
ORDER BY table_name;
```

Compare outputs; row counts should match for each table. Small differences here mean definite inconsistency, but equality does not guarantee identical data.

## 1.2 CHECKSUM TABLE on key tables

Pick a few important tables (e.g., accounts, orders).

On master:

```
USE repl_lab;  
CHECKSUM TABLE accounts, orders;
```

On slave:

```
USE repl_lab;  
CHECKSUM TABLE accounts, orders;  
Compare Checksum values per table; they must be identical.
```

## 2. Use Percona pt-table-checksum (recommended for full drift detection)

Install Percona Toolkit on the master host:

```
sudo apt update  
sudo apt install percona-toolkit  
pt-table-checksum --version
```

## 2.1 Run checksum from master (single command)

On master:

```
pt-table-checksum \
--host=master_host \
--user=root \
--password=YOURPASS \
--databases=repl_lab \
--replicate=percona.checksums \
--no-check-binlog-format \
--empty-replicate-table
```

What this does:

- Runs chunked checksum queries on master for each table in repl\_lab.
- Writes checksum for each chunk into percona.checksums on master.
- Those writes replicate to slaves, so percona.checksums is also created/updated there.

Sample output row (per table):

TS	ERRORS	DIFFS	ROWS	CHUNKS	SKIPPED	TIME	TABLE
12-27T11:10:01	0	0	1000	4	0	0.123	repl_lab.accounts
12-27T11:10:02	0	1	5000	10	0	0.456	repl_lab.orders

- DIFFS = 0 means no drift detected for that table.
- DIFFS > 0 means at least one chunk differs between master and slave(s).

## **2.2 Summarize differences per table**

On master (or any slave) query the checksum table:

**SELECT**

```
db,  
tbl,  
SUM(this_cnt) AS total_rows,  
SUM(master_cnt <> this_cnt OR master_crc <> this_crc  
OR ISNULL(master_crc) <> ISNULL(this_crc)) AS differing_chunks  
FROM percona.checksums  
GROUP BY db, tbl  
HAVING differing_chunks > 0;
```

If zero rows returned, all tables are consistent for all replicas.

## **3. Drill into which slave is inconsistent**

If you have multiple slaves, pt-table-checksum tracks per replica.

Show tables and replicas with differences:

```
pt-table-checksum \  
--replicate=percona.checksums \  
--replicate-check-only \  
--host=master_host \  
--user=root \  
--password=YOURPASS
```

This prints only tables with differences and identifies which slave(s) differ.

## **4. Lab steps to demonstrate consistency and drift**

### **4.1 Baseline: clean, consistent replication**

1. Ensure replication is running and healthy.

2. Run:

```
pt-table-checksum ... --databases=repl_lab
```

3. Confirm DIFFS=0 for all tables.

## 4.2 Simulate drift and detect it

On slave, make a bad manual change:

```
USE repl_lab;
```

```
UPDATE accounts SET balance = balance + 999 WHERE id = 1;
```

```
INSERT INTO accounts VALUES (999, 12345);
```

Run checksum again on master:

```
pt-table-checksum ... --databases=repl_lab
```

Now DIFFS for repl\_lab.accounts should be 1, and the percona.checksums table will show differing chunks.

Check specific differences:

```
SELECT db, tbl, chunk, this_cnt, master_cnt, this_crc, master_crc
```

```
FROM percona.checksums
```

```
WHERE db='repl_lab' AND tbl='accounts'
```

```
AND (master_cnt <> this_cnt OR master_crc <> this_crc);
```

## 5. (Optional) Auto-fix drift with pt-table-sync

Once you've identified drift, you can fix it by replaying changes from master to slave using pt-table-sync.

Example for one slave:

```
pt-table-sync \
--execute \
--replicate=percona.checksums \
--databases=repl_lab \
--host=master_host \
--user=root \
--password=YOURPASS

• Generates and executes INSERT/UPDATE/DELETE statements on master that replicate to the slave to fix the differences.

• Always test with --print first instead of --execute in real environments.
```

## 6. Minimal “no-tool” alternative (small lab)

If you don't want Percona Toolkit in a lab, you can do a simple checksum per table:

On master:

```
SELECT
'accounts' AS tbl,
COUNT(*) AS rows,
SUM(CRC32(CONCAT_WS('#', id, balance))) AS checksum
FROM repl_lab.accounts;
```

On slave, run the same and compare rows and checksum. Repeat for other tables. Crude but works for teaching purposes.

## 7. Lab verification checklist

- Row counts match between master and slave for all tables.
- CHECKSUM TABLE values match for key tables.
- pt-table-checksum run from master completes with no errors.
- percona.checksums shows DIFFS=0 for all tables when replication is healthy.
- After intentional drift on slave, pt-table-checksum flags the affected table.
- (Optional) pt-table-sync returns tables to a consistent state.

This gives a complete, hands-on procedure to verify and understand data consistency between master and slave, from quick checks to full tooling.